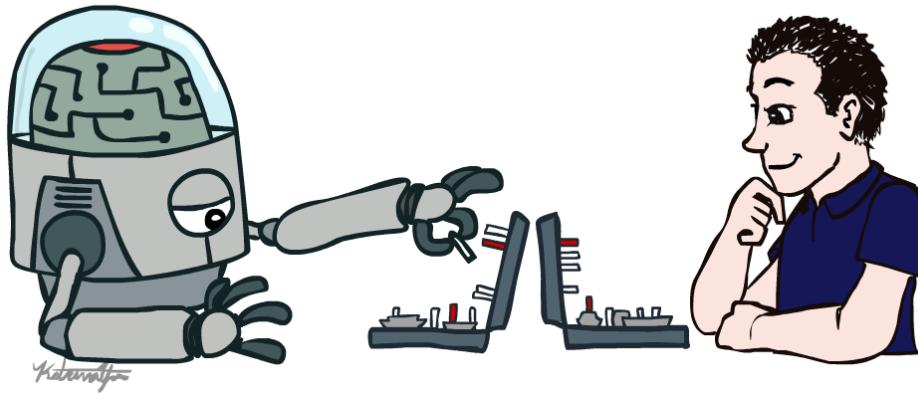


CS 4300/6300: Artificial Intelligence

Introduction



Instructor: Daniel Brown

University of Utah

[Adapted from those created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley. [http://ai.berkeley.edu.\]](http://ai.berkeley.edu.)

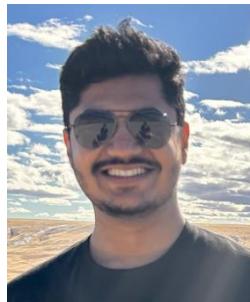
Course Staff

Professor

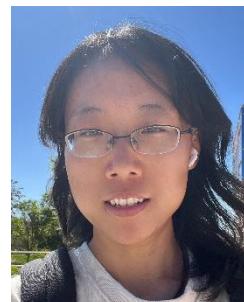


Daniel Brown

TAs



Atharv Belsare



Yile Li



Soumil Datta



Gabriella Goodman



Hannah Marumoto

Course Information

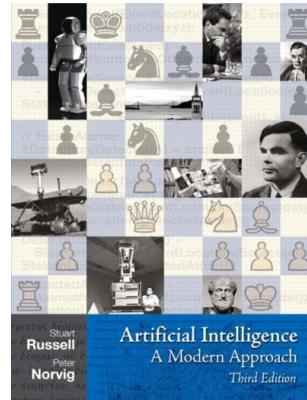
- **Communication:**
 - Announcements on Canvas
 - Questions? Discussion on piazza
- **Course format:**
 - Syllabus (see Canvas)
 - Homeworks and programming assignments turned in via Gradescope.
 - Homeworks must be completed individually.
 - Programming assignments can be done in pairs.
 - In class midterm and final.
- **Class Website:**
 - <https://dsbrown1331.github.io/intro-ai-class/>

Course Information

- There will be a lot of math (and programming)!
- Work and Grading:
 - 5 programming projects: Python,
 - 10% penalty for each day late (weekend counts as one day)
 - ~10 homework assignments.
 - Hands on experience working through math and algorithms.
 - Midterm (right before fall break) and Final (during finals week)
 - Class participation can help if your grade is on the margins.
 - We will occasionally have in class exercises/quizzes.
 - Pass/Fail based on whether you are in class or not.

Textbooks

- Russell & Norvig, AI: A Modern Approach, 3rd Ed.



Textbook

- Sutton and Barto “Reinforcement Learning: An Introduction”

<http://incompleteideas.net/book/ebook/the-book.html>

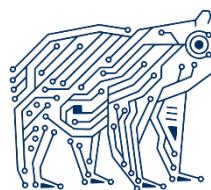
- Some notation in our class will be closer to Sutton and Barto than Russell and Norvig
- We will also sometimes have readings from
 - Deep Learning by Ian Goodfellow, Yoshua Bengio, and Aaroun Courville. MIT Press, 2016.
 - Decision Making Under Uncertainty by Mykel J. Kochenderfer. MIT Press, 2015.

Important This Week

- **Important this week:**
 - **Register** for the class on piazza --- our main resource for discussion and communication
 - **P0: Python crash course/tutorial** is out (due on Friday 8/29 by 11:59pm)
- **Also important:**
 - **Office Hours TBD**
 - **We will try and provide good coverage across days and times.**

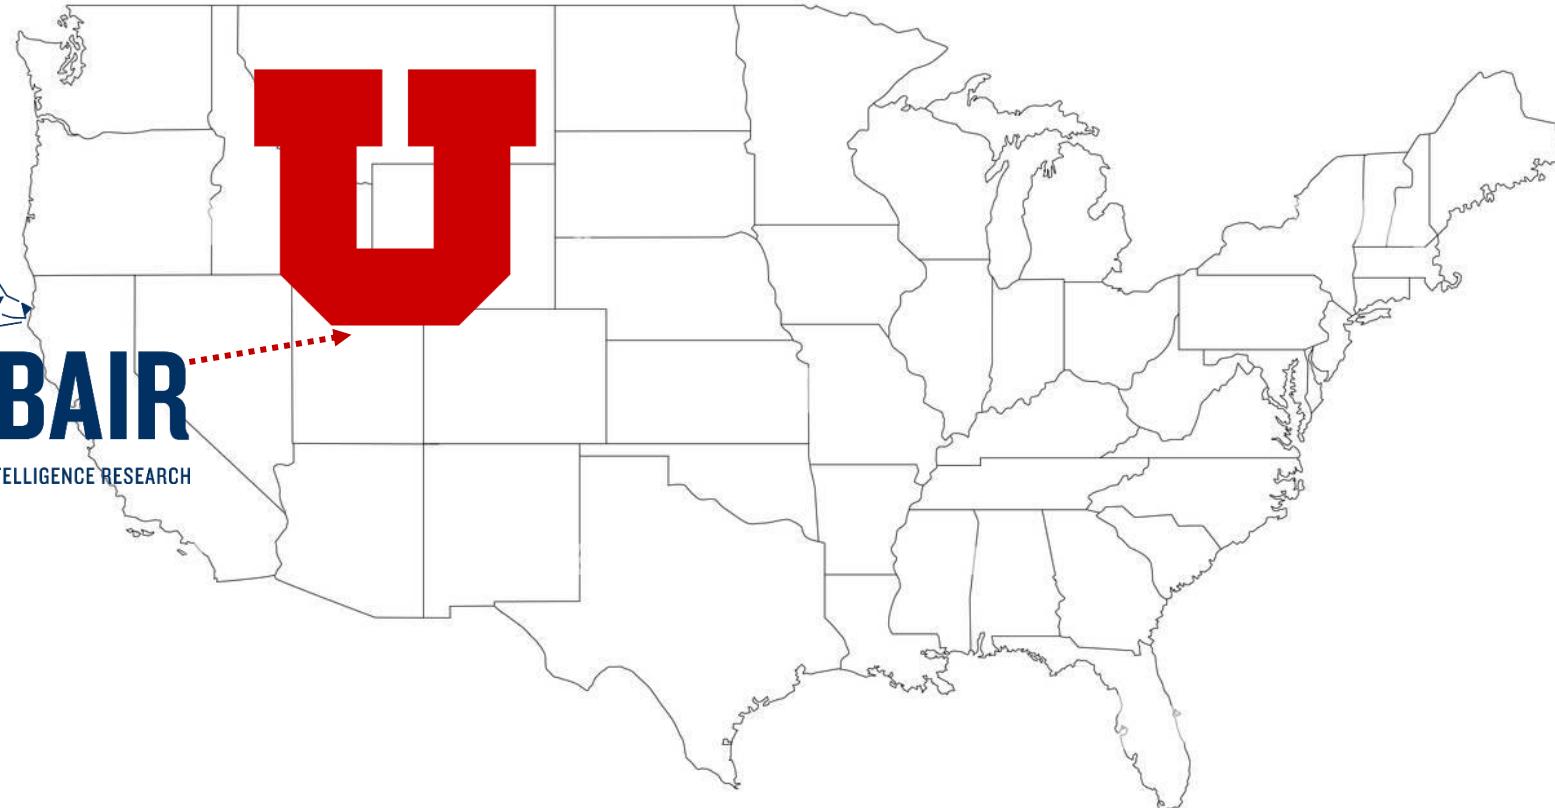
A little about me

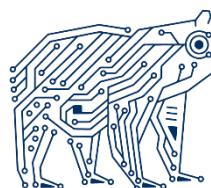




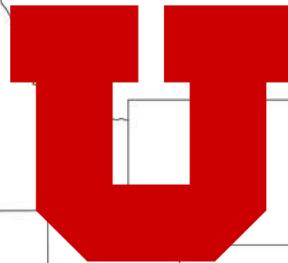
BAIR

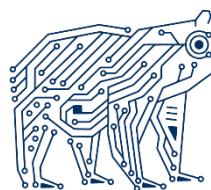
BERKELEY ARTIFICIAL INTELLIGENCE RESEARCH



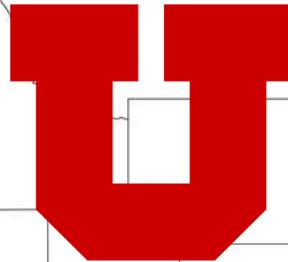


BAIR
BERKELEY ARTIFICIAL INTELLIGENCE RESEARCH

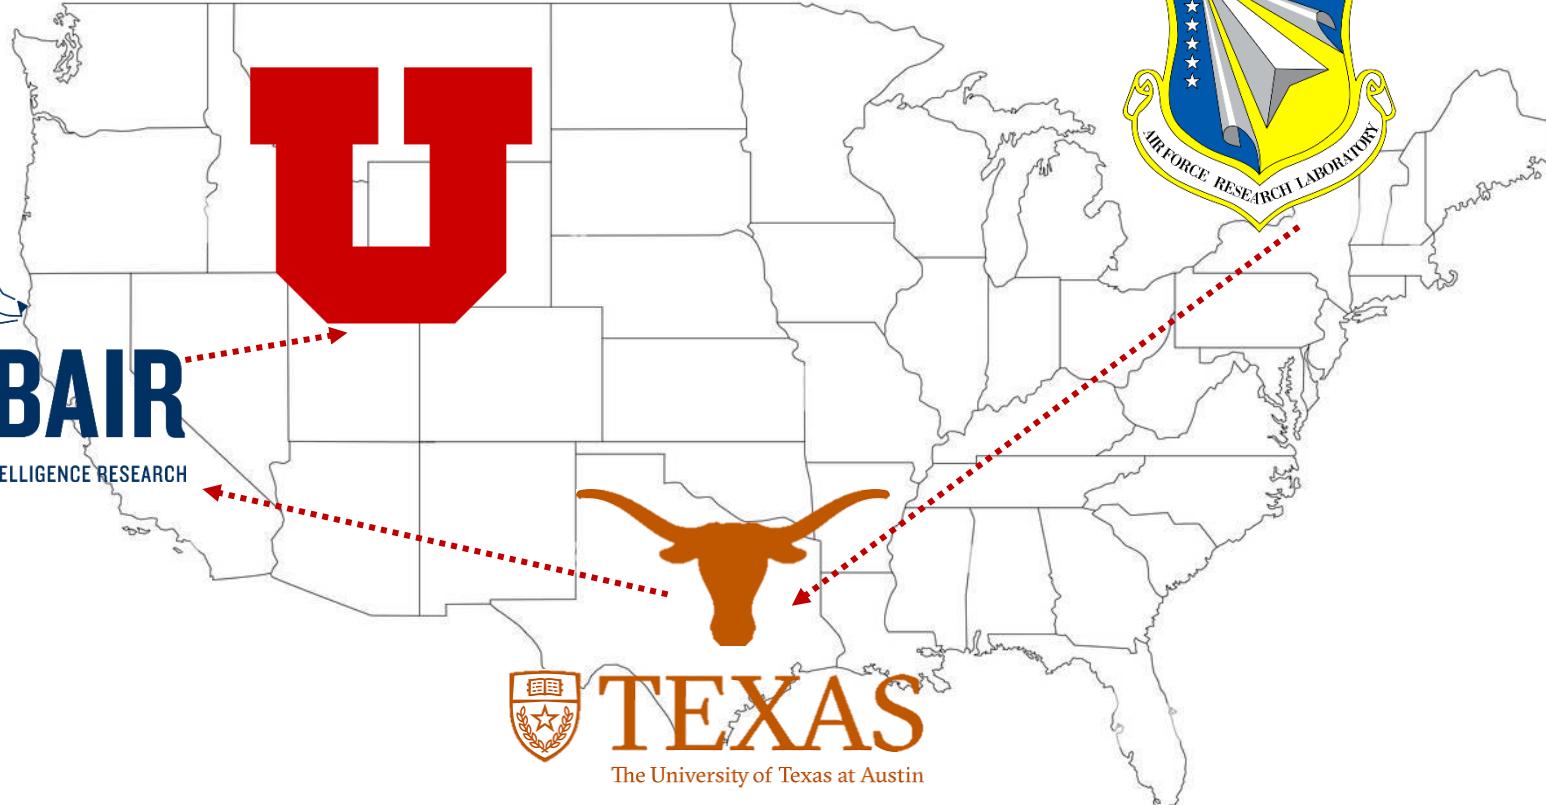


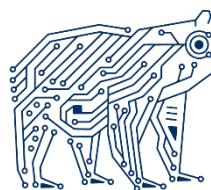


BAIR
BERKELEY ARTIFICIAL INTELLIGENCE RESEARCH



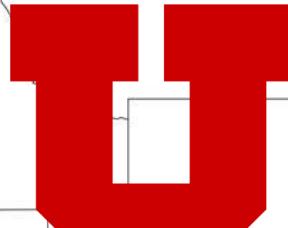
TEXAS
The University of Texas at Austin





BAIR

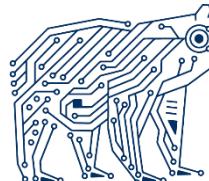
BERKELEY ARTIFICIAL INTELLIGENCE RESEARCH



TEXAS

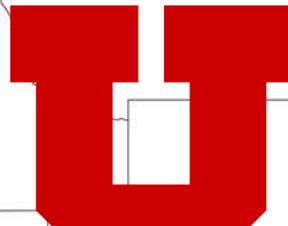
The University of Texas at Austin





BAIR

BERKELEY ARTIFICIAL INTELLIGENCE RESEARCH

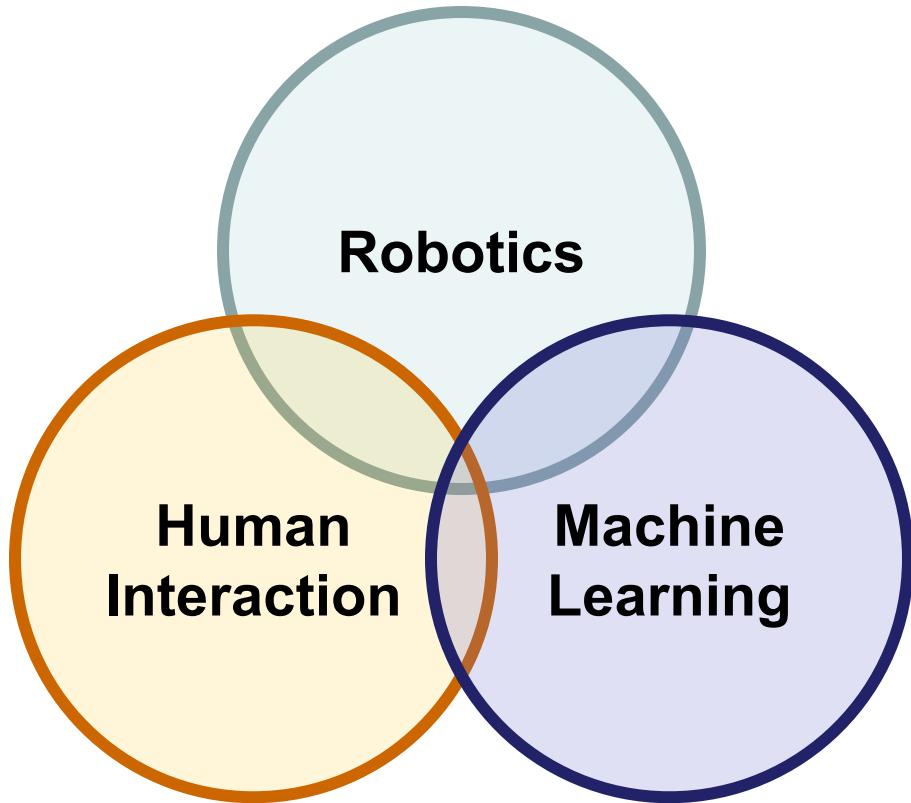


TEXAS

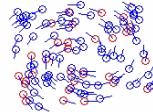
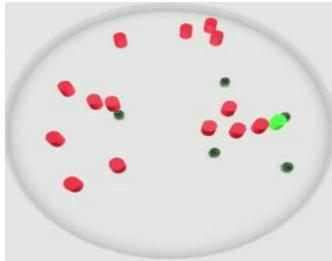
The University of Texas at Austin



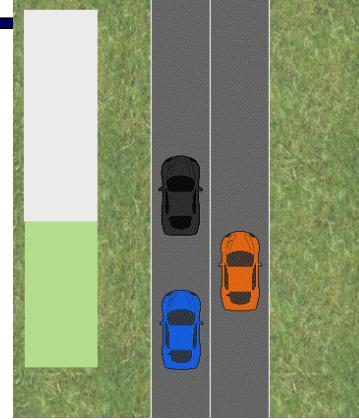




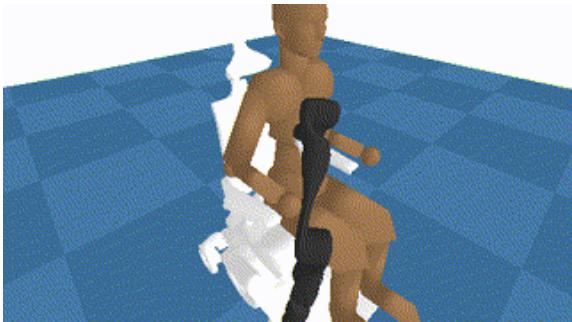
Human-Robot Interaction



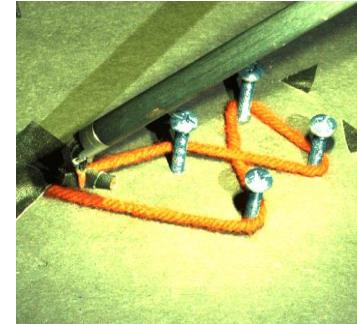
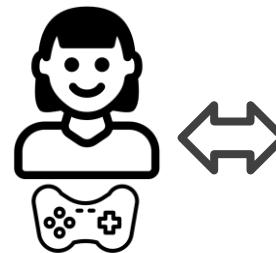
Human Swarm Interactions



Autonomous Driving

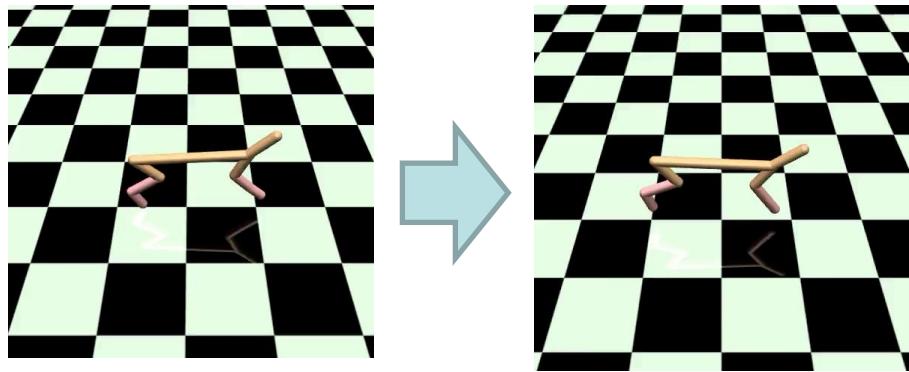


**Shared Autonomy and Assistive
Robotics**

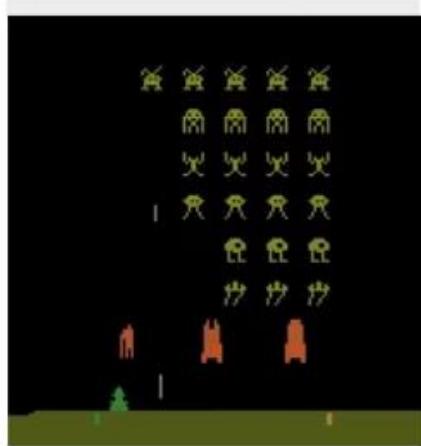


**Human-in-the-Loop
Robotics**

Learning models of human preferences



“Great weather today!”

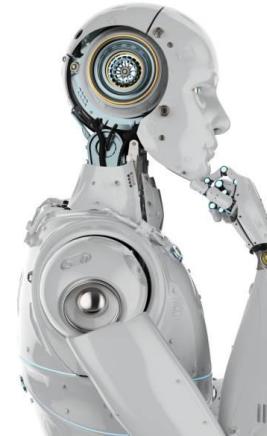


“Didn’t get the job offer...”

AI Safety and Robustness

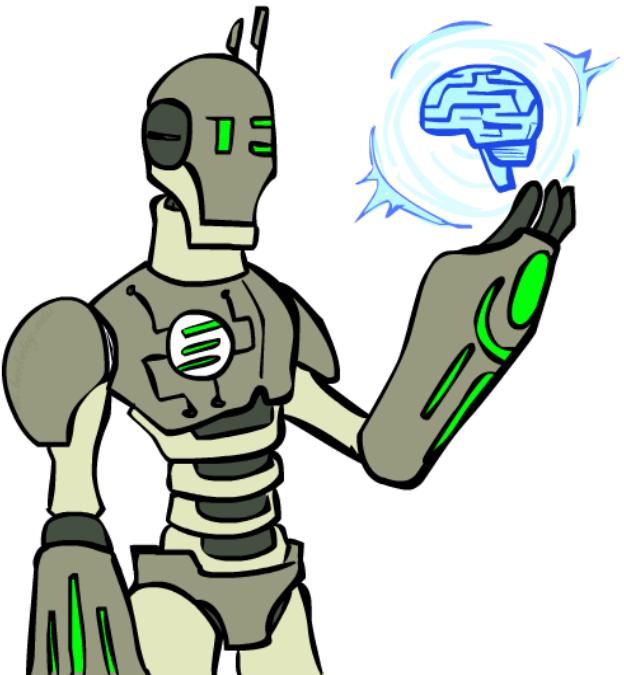


	UNCH	-47.20	47.28	94.45	61.73	13,773,000	W/2 New York
6.50	UNCH	-65.30	66.90	66.76	66.05	7,701,000	W/2 New York
7.03	-82	85.00	87.95	88.54	87.01	4,307,500	W/2 New York
39.16	-85	87.53	41.14	39.70	39.13	2,751,000	W/2 New York
23.18	+.17	23.20	23.25	23.54	22.64	9,495,000	W/2 New York
16.71	-.06	16.71	16.74	16.89	16.53	37,306,407	W/2 New York
35.39	-.18	35.00	36.35	35.87	35.34	8,059,000	W/2 New York
69.47	-.65	69.25	70.49	70.10	69.34	3,676,000	W/2 New York
24.41	-.73	24.37	24.65	25.10	24.41	4,468,000	W/2 New York
49.24	-.45	48.50	50.32	49.87	49.24	1,729,000	W/2 New York
42.58	-.36	42.05	44.25	43.02	42.47	9,740,500	W/2 New York
63.16	-.52	62.78	64.14	63.71	62.70	9,275,000	W/2 New York
83.48	-.86	83.07	84.00	84.30	83.12	10,236,400	W/2 New York

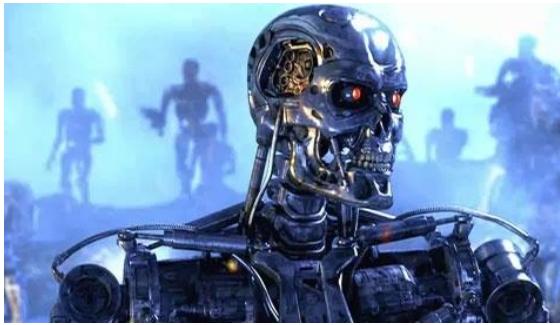


Today

- What is artificial intelligence?
- What can AI do?
- What is this course?



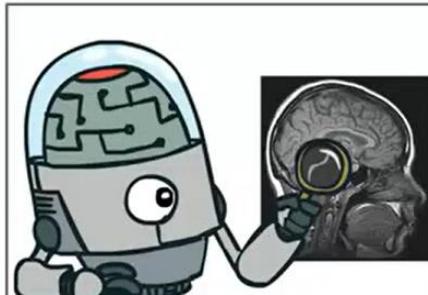
Sci-Fi AI?



What is AI?

The science of making machines that:

Think like people



Think rationally



Act like people



Act rationally



Rational Decisions

We'll use the term **rational** in a very specific, technical way:

- Rational: maximally achieving pre-defined goals
- Rationality only concerns what decisions are made
(not the thought process behind them)
- Goals are expressed in terms of the **utility** of outcomes
- Being rational means **maximizing your expected utility**

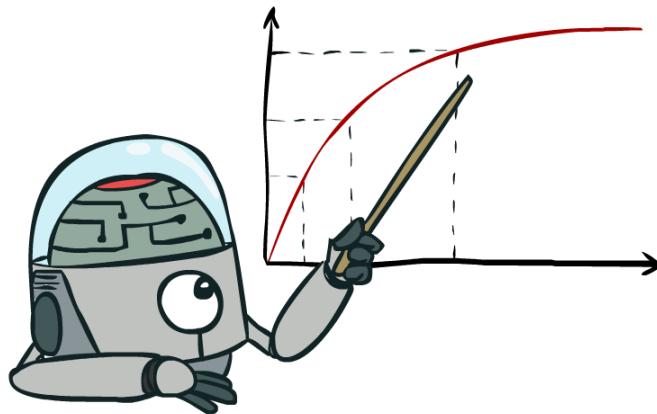
One possible title for this course would be:

Computational Rationality

Another Possible Course Title

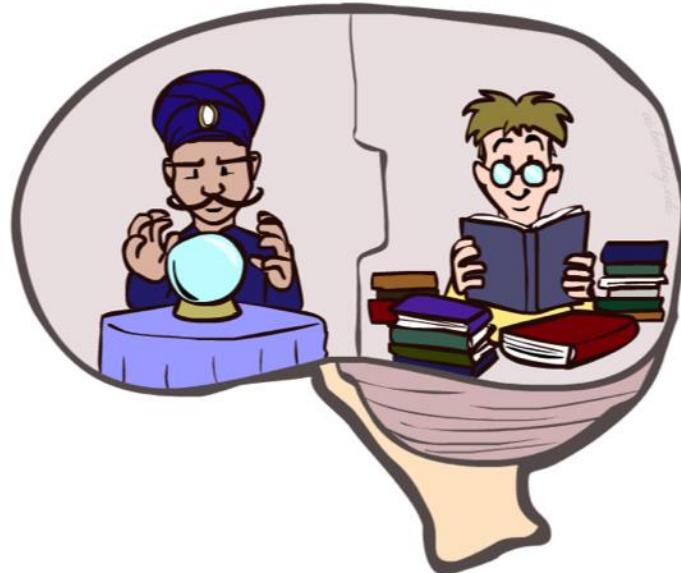
- Foundational Mathematics and Algorithms for Sequential Decision Making Under Uncertainty!

Maximize Your Expected Utility

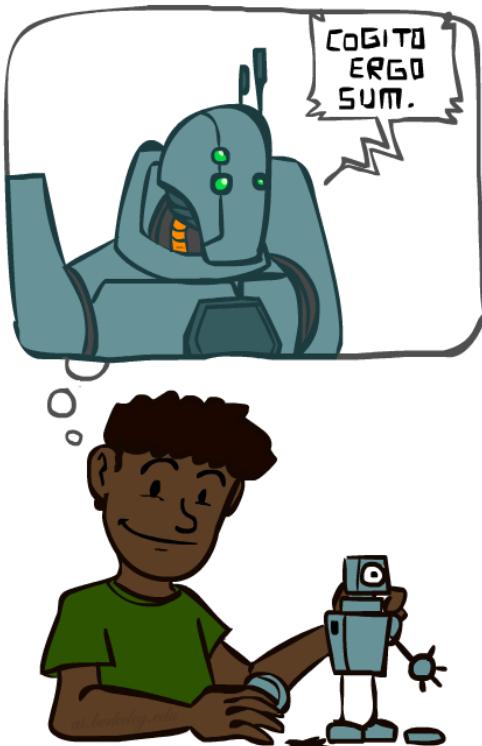


What About the Brain?

- Brains (human minds) are very good at making rational decisions, but not perfect
- Brains aren't as modular as software, so hard to reverse engineer!
- "Brains are to intelligence as wings are to flight"
- Lessons learned from the brain: memory and simulation are key to decision making



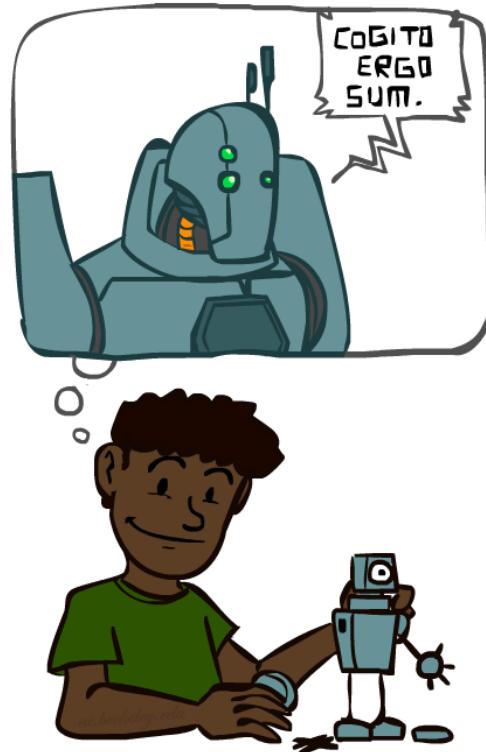
A (Short) History of AI





A (Short) History of AI

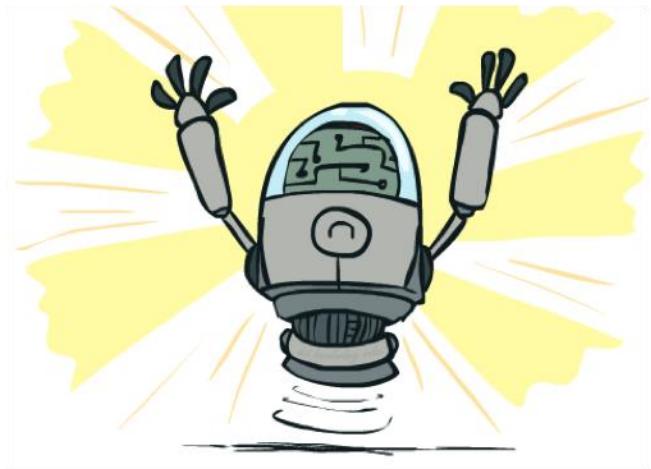
- 1940–1950: Early days
 - 1943: McCulloch & Pitts: Boolean circuit model of brain
 - 1950: Turing's "Computing Machinery and Intelligence"
- 1950–70: Excitement: Look, Ma, no hands!
 - 1950s: Early AI programs, including Samuel's checkers program, Newell & Simon's Logic Theorist, Gelernter's Geometry Engine
 - 1956: Dartmouth meeting: "Artificial Intelligence" adopted
 - 1965: Robinson's complete algorithm for logical reasoning
- 1970–90: Knowledge-based approaches
 - 1969–79: Early development of knowledge-based systems
 - 1980–88: Expert systems industry booms
 - 1988–93: Expert systems industry busts: "AI Winter"
- 1990–: Statistical approaches
 - Resurgence of probability and statistics, focus on uncertainty
 - General increase in technical depth
 - Agents and learning systems... "AI Spring"?
- 2014–: Deep Learning and Scaling Things Up
 - Big data
 - Powerful compute (e.g. GPUs)
 - ChatGPT and other Large Language Models (LLMs)



What Can AI Do?

Quiz: Which of the following can be done at present?

- Play a decent game of table tennis?
- Play a decent game of Jeopardy?
- Drive safely along a curving mountain road?
- Buy a week's worth of groceries on the web?
- Buy a week's worth of groceries at Walmart?
- Discover and prove a new mathematical theorem?
- Converse successfully with another person for an hour?
- Perform a surgical operation?
- Come into your apartment and unload your dishwasher
- Translate spoken Chinese into spoken English in real time?
- Write decent poetry?
- Create aesthetically pleasing artwork
- Fold your laundry?



Natural Language

- Speech technologies (e.g. Siri)
 - Automatic speech recognition (ASR)
 - Text-to-speech synthesis (TTS)
 - Dialog systems
- Language processing technologies
 - Question answering
 - Machine translation

"Il est impossible aux journalistes de rentrer dans les régions tibétaines"

Bruno Philip, correspondant du "Monde" en Chine, estime que les journalistes de l'AFP qui ont été expulsés de la province tibétaine du Qinghai "n'étaient pas dans l'ilégalité".

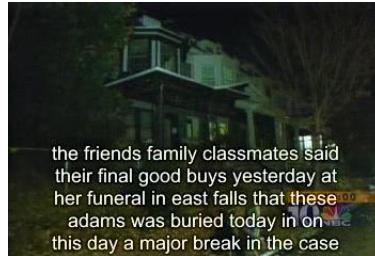
Les faits Le dalaï-lama dénonce l'"enfer" imposé au Tibet depuis sa fuite, en 1959
Vidéo Anniversaire de la rébellion



"It is impossible for journalists to enter Tibetan areas"

Philip Bruno, correspondent for "World" in China, said that journalists of the AFP who have been deported from the Tibetan province of Qinghai "were not illegal."

Facts The Dalai Lama denounces the "hell" imposed since he fled Tibet in 1959
Video Anniversary of the Tibetan rebellion: China on guard



- Web search
- Text classification, spam filtering, etc...

Chat GPT

chat gpt - Search

ChatGPT: Optimizing Lang

404 Not Found

chat gpt 404 Not Found -

Did we just crash ChatGPT?

Not syncing

Introducing ChatGPT research release Try Learn more >

OpenAI

API RESEARCH BLOG ABOUT

ChatGPT: Optimizing Language Models for Dialogue

We've trained a model called ChatGPT which interacts in a conversational way. The dialogue format makes it possible for ChatGPT to answer followup questions, admit its mistakes, challenge incorrect premises, and reject inappropriate requests. ChatGPT is a sibling model to InstructGPT, which is trained to follow an instruction in a prompt and provide a detailed response.

TRY CHATGPT ↗

43°F Raining now

Search

3:22 PM 10
1/10/2023



Text to Images (DALL-E)

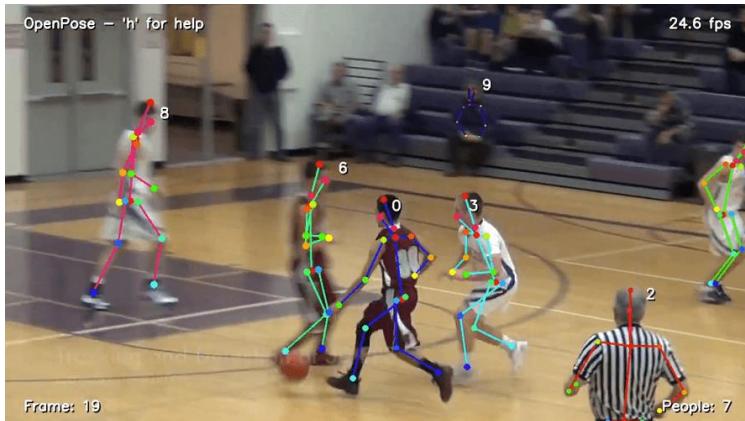
- “An astronaut riding a horse in a photo-realistic way”



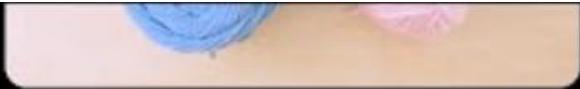
- “An armchair in the shape of an avocado.”

Vision (Perception)

- Object and face recognition
- Scene segmentation
- Image classification

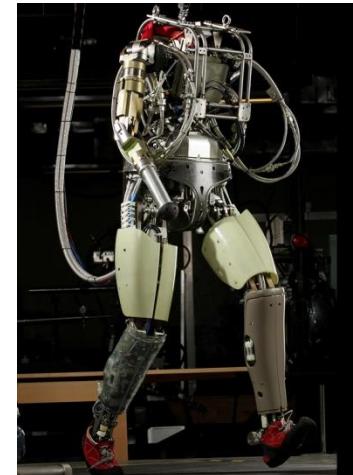
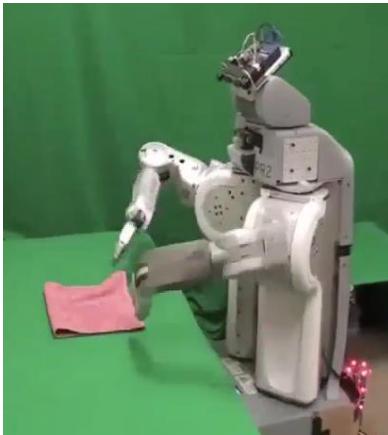


Multi-Modal Models



Robotics

- Robotics
 - Part mech. eng.
 - Part AI
 - Reality much harder than simulations!
- Technologies
 - Autonomous Vehicles
 - Rescue
 - Soccer!
 - Lots of automation...
 - But increasing amounts of learning...
- In this class:
 - We ignore mechanical aspects
 - Methods for planning
 - Methods for control



Images from UC Berkeley, Boston Dynamics, RoboCup, Google

Boston Dynamics Atlas



<https://vision-locomotion.github.io/>



Do As I Can, Not As I Say: Grounding Language in Robotic Affordances

¹ Michael Ahn*, Anthony Brohan*, Noah Brown*, Yevgen Chebotar*, Omar Cortes*, Byron David*, Chelsea Finn*, Chuyuan Fu[†], Keerthana Gopalakrishnan*, Karol Hausman*, Alex Herzog[†], Daniel Ho[†], Jasmine Hsu*, Julian Ibarz*, Brian Ichter*, Alex Irpan*, Eric Jang*, Rosario Jauregui Ruano*, Kyle Jeffrey*, Sally Jesmonth*, Nikhil J Joshi*, Ryan Julian*, Dmitry Kalashnikov*, Yuheng Kuang*, Kuang-Huei Lee*, Sergey Levine*, Yao Lu*, Linda Luu*, Carolina Parada*, Peter Pastor[†], Jornell Quiambao*, Kanishka Rao*, Jarek Rettinghouse*, Diego Reyes*, Pierre Sermanet*, Nicolas Sievers*, Clayton Tan*, Alexander Toshev*, Vincent Vanhoucke*, Fei Xia*, Ted Xiao*, Peng Xu*, Sichun Xu*, Mengyuan Yan[†], Andy Zeng*

*Robotics at Google, [†]Everyday Robots





autonomous, 1x speed

π

Logic

- Logical systems
 - Theorem provers
 - NASA fault diagnosis
 - Question answering
- Methods:
 - Deduction systems
 - Constraint satisfaction
 - Satisfiability solvers (huge advances!)

Article

Discovering faster matrix multiplication algorithms with reinforcement learning

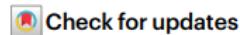
<https://doi.org/10.1038/s41586-022-05172-4>

Received: 2 October 2021

Accepted: 2 August 2022

Published online: 5 October 2022

Open access



Alhussein Fawzi^{1,2}✉, Matej Balog^{1,2}, Aja Huang^{1,2}, Thomas Hubert^{1,2}, Bernardino Romera-Paredes^{1,2}, Mohammadamin Barekatain¹, Alexander Novikov¹, Francisco J. R. Ruiz¹, Julian Schrittwieser¹, Grzegorz Swirszcz¹, David Silver¹, Demis Hassabis¹ & Pushmeet Kohli¹

Improving the efficiency of algorithms for fundamental computations can have a widespread impact, as it can affect the overall speed of a large amount of computations. Matrix multiplication is one such primitive task, occurring in many systems—from neural networks to scientific computing routines. The automatic discovery of algorithms using machine learning offers the prospect of reaching beyond human intuition and outperforming the current best human-designed algorithms. However, automating the algorithm discovery procedure is intricate, as the space of possible algorithms is enormous. Here we report a deep reinforcement learning approach based on AlphaZero¹ for discovering efficient and provably correct algorithms for the multiplication of arbitrary matrices. Our agent, AlphaTensor, is trained to play a single-player game where the objective is finding tensor decompositions within a finite factor space. AlphaTensor discovered algorithms that outperform the state-

SCIENCE

AlphaGeometry: An Olympiad-level AI system for geometry

17 JANUARY 2024

Trieu Trinh and Thang Luong

RESEARCH

Advanced version of Gemini with Deep Think officially achieves gold-medal standard at the International Mathematical Olympiad

21 JULY 2025

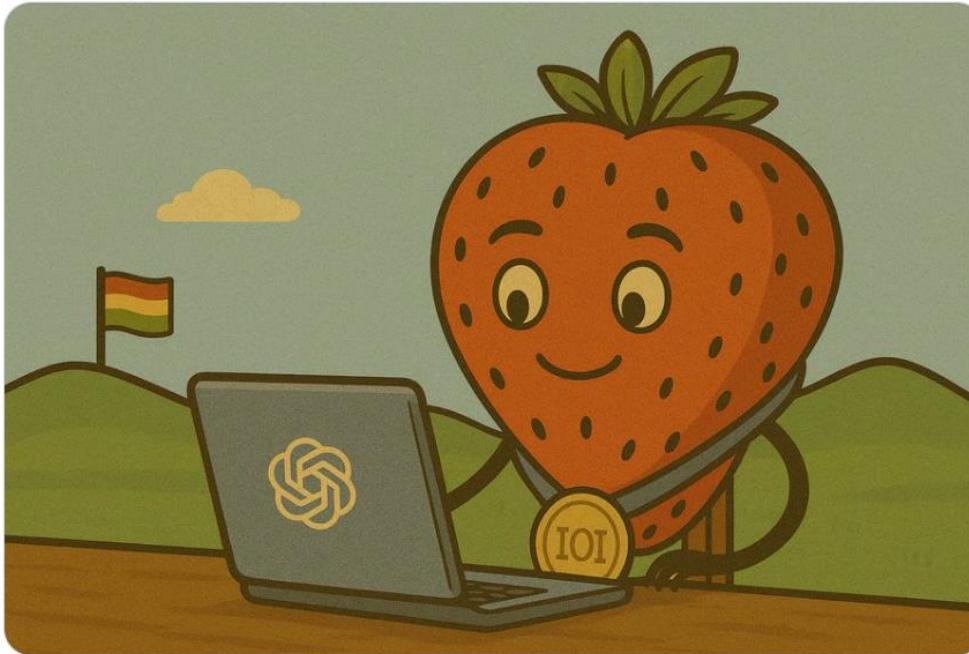
Thana Luona and Edward Lockhart



Sheryl Hsu
@SherylHsu02



1/n I'm thrilled to share that our @OpenAI reasoning system scored high enough to achieve gold 🥇🥇 in one of the world's top programming competitions - the 2025 International Olympiad in Informatics (IOI) - placing first among AI participants! 🤖💻



12:00 PM · Aug 11, 2025 · 2.2M Views

Game Playing

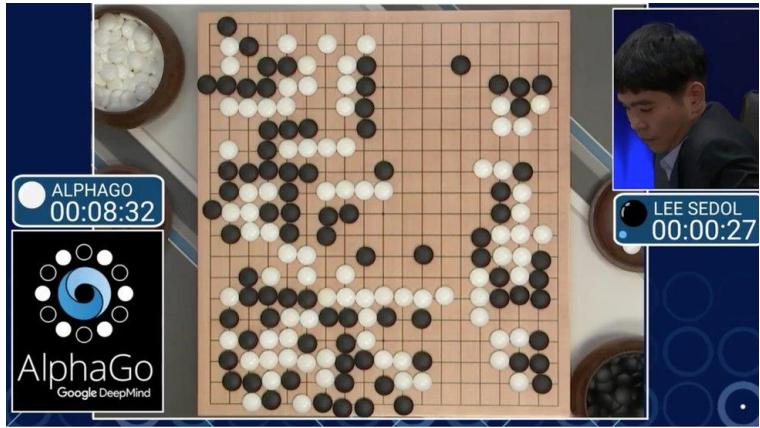
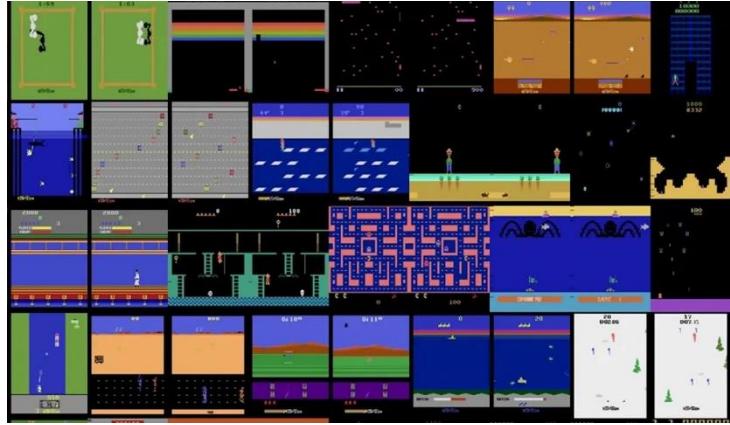
- **Classic Moment: May, '97: Deep Blue vs. Kasparov**
 - First match won against world champion
 - “Intelligent creative” play
 - 200 million board positions per second
 - Humans understood 99.9 of Deep Blue's moves
 - Can do about the same now with a PC cluster
- **Open question:**
 - How does human cognition deal with the search space explosion of chess?
 - Or: how can humans compete with computers at all??
- **1996: Kasparov Beats Deep Blue**

“I could feel --- I could smell --- a new kind of intelligence across the table.”
- **1997: Deep Blue Beats Kasparov**

“Deep Blue hasn't proven anything.”
- **Huge game-playing advances recently, e.g. in Go!**



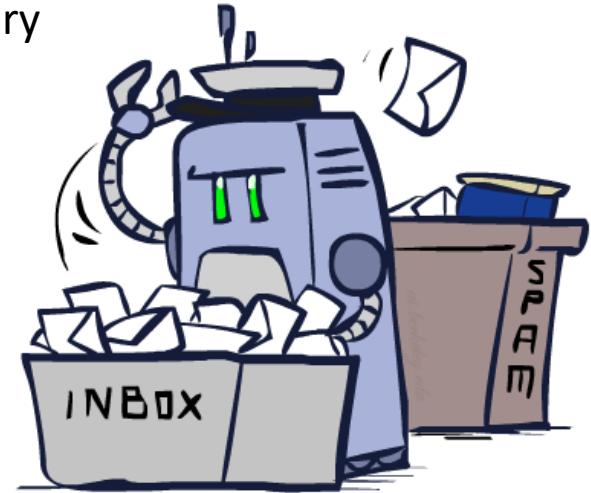
Winning at Games



Decision Making

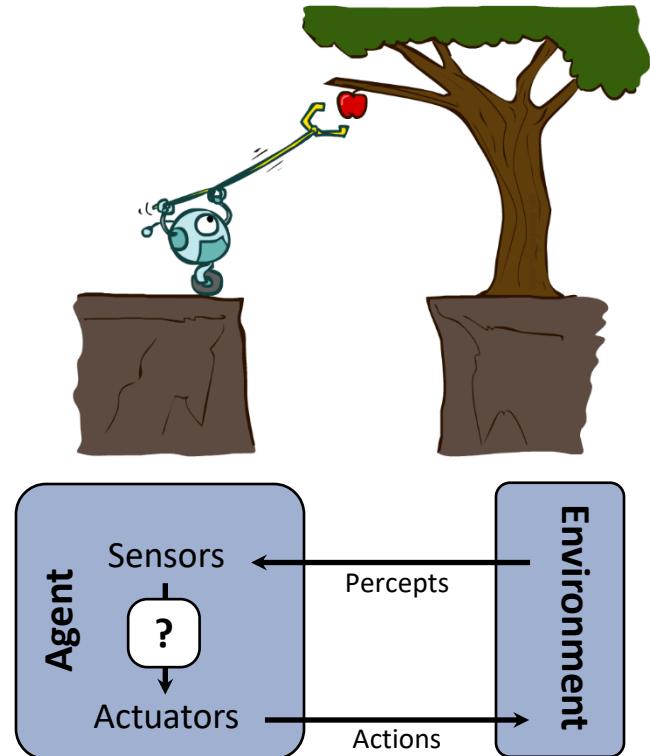
- Applied AI involves many kinds of automation

- Scheduling, e.g. airline routing, military
- Route planning, e.g. Google maps
- Medical diagnosis
- Web search engines
- Spam classifiers
- Automated help desks
- Fraud detection
- Product recommendations
- ... Lots more!

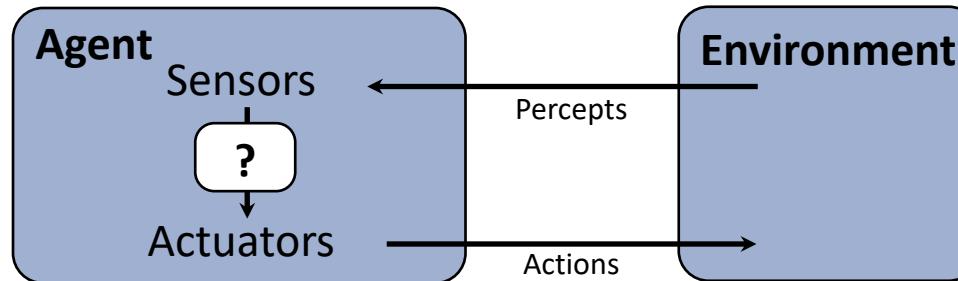
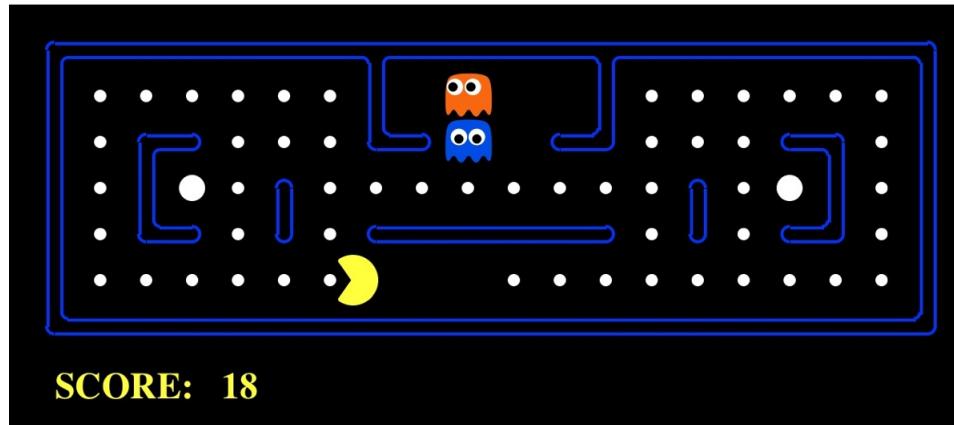


Designing Rational Agents

- An **agent** is an entity that *perceives* and *acts*.
- A **rational agent** selects actions that maximize its (expected) **utility**.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions
- **This course** is about:
 - General AI techniques for a variety of problem types
 - Learning to recognize when and how a new problem can be solved with an existing technique



Pac-Man as an Agent



Pacman

- Google: “play pacman doodle”

Course Topics

- Part I: Making Decisions
 - Fast search
 - Adversarial and uncertain search
- Part II: Markov Decision Process (MDP)
 - Planning
 - Supervised Learning
 - Reinforcement learning
- Part III: Reasoning Under Uncertainty
 - Bayes Nets
 - Hidden Markov Models (HMMs)
 - Imitation Learning
 - Reinforcement Learning from Human Feedback

